

Logistic Regression

Goal:

The goal of this assignment is to train various Logistic Regression models on KDDCup99 dataset. You can find more details and download the dataset here: <http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html>

Tasks:

1. Become familiar with Scikit-Learn's logistic regression models. You can find more details and examples here: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html
2. Download KDD99.csv dataset (<http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html>)
3. Using the Scikit-Learn Library, train the Logistic Regression model to classify Probe vs Not Probe classes using all available features. Probe class consists of all probing attacks, such as ipsweep probe, nmap probe, portsweep probe, etc. Non probe class consists of every other attack in the dataset. Make sure you split your data into train and test sets. In general, 80% for training 20% for testing is a good split.
4. Summarize your results. What is the training accuracy? What is the testing accuracy? Do you think your models overfits? How many iterations did it take to converge?
5. Take a look at the trained model parameters. Which features have the largest weights in the absolute value? These are potentially most informative features for the trained model. Do they make sense? Why?
6. Train model with L1 and L2 regularization. Compare the accuracies for L1 and L2 and plot the feature weights.
7. Try different strengths for L2 regularization. (You can use values .001, .01, .1, 1, 10 etc.) Show how the train and test accuracy varies with different regularization strengths.
8. Carry out all the tasks on a single Jupyter notebook. Discuss your findings when appropriate. Is what you observe expected? Why?

Installing the required libraries

```
In [1]: !pip3 install plotly
```

Requirement already satisfied: plotly in ./opt/anaconda3/lib/python3.8/site-packages (5.3.1)
Requirement already satisfied: six in ./opt/anaconda3/lib/python3.8/site-packages (from plotly) (1.15.0)
Requirement already satisfied: tenacity>=6.2.0 in ./opt/anaconda3/lib/python3.8/site-packages (from plotly) (8.0.1)

```
In [2]: !pip3 install seaborn
```

Requirement already satisfied: seaborn in ./opt/anaconda3/lib/python3.8/site-packages (0.11.1)
Requirement already satisfied: numpy>=1.15 in ./opt/anaconda3/lib/python3.8/site-packages (from seaborn) (1.20.1)
Requirement already satisfied: scipy>=1.0 in ./opt/anaconda3/lib/python3.8/site-packages (from seaborn) (1.6.2)
Requirement already satisfied: pandas>=0.23 in ./opt/anaconda3/lib/python3.8/site-packages (from seaborn) (1.2.4)
Requirement already satisfied: matplotlib>=2.2 in ./opt/anaconda3/lib/python3.8/site-packages (from seaborn) (3.3.4)
Requirement already satisfied: pillow>=6.2.0 in ./opt/anaconda3/lib/python3.8/site-packages (from matplotlib>=2.2->seaborn) (8.2.0)
Requirement already satisfied: python-dateutil>=2.1 in ./opt/anaconda3/lib/python3.8/site-packages (from matplotlib>=2.2->seaborn) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in ./opt/anaconda3/lib/python3.8/site-packages (from matplotlib>=2.2->seaborn) (1.3.1)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in ./opt/anaconda3/lib/python3.8/site-packages (from matplotlib>=2.2->seaborn) (2.4.7)
Requirement already satisfied: cycler>=0.10 in ./opt/anaconda3/lib/python3.8/site-packages (from matplotlib>=2.2->seaborn) (0.10.0)
Requirement already satisfied: six in ./opt/anaconda3/lib/python3.8/site-packages (from cyclar>=0.10->matplotlib>=2.2->seaborn) (1.15.0)
Requirement already satisfied: pytz>=2017.3 in ./opt/anaconda3/lib/python3.8/site-packages (from pandas>=0.23->seaborn) (2021.1)

```
In [3]: import plotly.express as px
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

Datafiles

```
In [4]: df = pd.read_csv('kddcup99.csv')
df
```

Out[4]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	...	dst_host_srv_count	dst_host_same_srv_rate	dst_h
0	0	tcp	http	SF	181	5450	0	0	0	0	...	9	1.0	
1	0	tcp	http	SF	239	486	0	0	0	0	...	19	1.0	
2	0	tcp	http	SF	235	1337	0	0	0	0	...	29	1.0	
3	0	tcp	http	SF	219	1337	0	0	0	0	...	39	1.0	
4	0	tcp	http	SF	217	2032	0	0	0	0	...	49	1.0	
...	
494015	0	tcp	http	SF	310	1881	0	0	0	0	...	255	1.0	
494016	0	tcp	http	SF	282	2286	0	0	0	0	...	255	1.0	
494017	0	tcp	http	SF	203	1200	0	0	0	0	...	255	1.0	
494018	0	tcp	http	SF	291	1200	0	0	0	0	...	255	1.0	
494019	0	tcp	http	SF	219	1234	0	0	0	0	...	255	1.0	

494020 rows x 42 columns

```
In [5]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 494020 entries, 0 to 494019
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   duration                              494020 non-null  int64
1   protocol_type                         494020 non-null  object
2   service                              494020 non-null  object
3   flag                                  494020 non-null  object
4   src_bytes                             494020 non-null  int64
5   dst_bytes                             494020 non-null  int64
6   land                                  494020 non-null  int64
7   wrong_fragment                        494020 non-null  int64
8   urgent                                494020 non-null  int64
9   hot                                    494020 non-null  int64
10  num_failed_logins                     494020 non-null  int64
11  logged_in                             494020 non-null  int64
12  lnum_compromised                      494020 non-null  int64
13  lroot_shell                           494020 non-null  int64
14  lsu_attempted                         494020 non-null  int64
15  lnum_root                             494020 non-null  int64
16  lnum_file_creations                   494020 non-null  int64
17  lnum_shells                           494020 non-null  int64
18  lnum_access_files                     494020 non-null  int64
19  lnum_outbound_cmds                   494020 non-null  int64
20  is_host_login                         494020 non-null  int64
21  is_guest_login                        494020 non-null  int64
22  count                                 494020 non-null  int64
23  srv_count                             494020 non-null  int64
24  serror_rate                           494020 non-null  float64
25  srv_serror_rate                       494020 non-null  float64
26  rerror_rate                           494020 non-null  float64
27  srv_rerror_rate                       494020 non-null  float64
28  same_srv_rate                         494020 non-null  float64
29  diff_srv_rate                         494020 non-null  float64
30  srv_diff_host_rate                   494020 non-null  float64
31  dst_host_count                        494020 non-null  int64
32  dst_host_srv_count                   494020 non-null  int64

-- -- --
33  dst_host_same_srv_rate               494020 non-null  float64
34  dst_host_diff_srv_rate               494020 non-null  float64
35  dst_host_same_src_port_rate          494020 non-null  float64
36  dst_host_srv_diff_host_rate          494020 non-null  float64
37  dst_host_serror_rate                 494020 non-null  float64
38  dst_host_srv_serror_rate             494020 non-null  float64
39  dst_host_rerror_rate                 494020 non-null  float64
40  dst_host_srv_rerror_rate             494020 non-null  float64
41  label                                494020 non-null  object
dtypes: float64(15), int64(23), object(4)
memory usage: 158.3+ MB
```

Encoding string

```
In [6]: df['label'].unique()

Out[6]: array(['normal', 'buffer_overflow', 'loadmodule', 'perl', 'neptune',
              'smurf', 'guess_passwd', 'pod', 'teardrop', 'portsweep', 'ipsweep',
              'land', 'ftp_write', 'back', 'imap', 'satan', 'phf', 'nmap',
              'multihop', 'warezmaster', 'warezclient', 'spy', 'rootkit'],
              dtype=object)

In [7]: df['label']=df['label'].replace(['ipsweep','nmap','portsweep'],1)
df['label']=df['label'].replace(['normal','buffer_overflow','loadmodule','perl','neptune','smurf','guess_passwd','po
```

```
In [8]: df
```

Out[8]:

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	...	dst_host_srv_count	dst_host_same_srv_rate	dst_h
0	0	tcp	http	SF	181	5450	0	0	0	0	...	9	1.0	
1	0	tcp	http	SF	239	486	0	0	0	0	...	19	1.0	
2	0	tcp	http	SF	235	1337	0	0	0	0	...	29	1.0	
3	0	tcp	http	SF	219	1337	0	0	0	0	...	39	1.0	
4	0	tcp	http	SF	217	2032	0	0	0	0	...	49	1.0	
...	
494015	0	tcp	http	SF	310	1881	0	0	0	0	...	255	1.0	
494016	0	tcp	http	SF	282	2286	0	0	0	0	...	255	1.0	
494017	0	tcp	http	SF	203	1200	0	0	0	0	...	255	1.0	
494018	0	tcp	http	SF	291	1200	0	0	0	0	...	255	1.0	
494019	0	tcp	http	SF	219	1234	0	0	0	0	...	255	1.0	

```
In [9]: dummies1=pd.get_dummies(df.protocol_type)
dummies1.head()
```

```
Out[9]:
```

	icmp	tcp	udp
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0

```
In [10]: dummies2=pd.get_dummies(df.service)
dummies2.head()
```

```
Out[10]:
```

	IRC	X11	Z39_50	auth	bgp	courier	csnet_ns	ctf	daytime	discard	...	telnet	tftp_u	tim_i	time	urh_i	urp_i	uucp	uucp_path	vmnet	whois
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows x 66 columns

```
In [11]: dummies3=pd.get_dummies(df.flag)
dummies3.head()
```

```
Out[11]:
```

	OTH	REJ	RSTO	RSTOS0	RSTR	S0	S1	S2	S3	SF	SH
0	0	0	0	0	0	0	0	0	0	1	0
1	0	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	1	0
3	0	0	0	0	0	0	0	0	0	1	0
4	0	0	0	0	0	0	0	0	0	1	0

```
In [12]: df1=pd.concat([df,dummies1,dummies2,dummies3],axis='columns')
df1.drop(['protocol_type','service','flag'], axis='columns',inplace= True)
df1.head(5)
```

```
Out[12]:
```

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	num_failed_logins	logged_in	lnum_compromised	...	REJ	RSTO	RSTOS0	RSTR	S0
0	0	181	5450	0	0	0	0	0	1	0	...	0	0	0	0	0
1	0	239	486	0	0	0	0	0	1	0	...	0	0	0	0	0
2	0	235	1337	0	0	0	0	0	1	0	...	0	0	0	0	0
3	0	219	1337	0	0	0	0	0	1	0	...	0	0	0	0	0
4	0	217	2032	0	0	0	0	0	1	0	...	0	0	0	0	0

5 rows x 119 columns

```
In [13]: df1.isnull().sum()
```

```
Out[13]: duration      0
src_bytes      0
dst_bytes      0
land      0
wrong_fragment  0
..
S1      0
S2      0
S3      0
SF      0
SH      0
Length: 119, dtype: int64
```

```
In [14]: from sklearn.preprocessing import MinMaxScaler
```

```
In [15]: numeric_cols = df1.select_dtypes(include=np.number).columns.tolist()
categorical_cols = df1.select_dtypes('object').columns.tolist()
categorical_cols
```

```
Out[15]: []
```

```
In [16]: scaler = MinMaxScaler()
```

```
In [17]: scaler.fit(df1[numeric_cols])
```

```
Out[17]: MinMaxScaler()
```

```
In [18]: df1[numeric_cols] = scaler.transform(df1[numeric_cols])
df1[numeric_cols]
```

Out[18]:

	duration	src_bytes	dst_bytes	land	wrong_fragment	urgent	hot	num_failed_logins	logged_in	Inum_compromised	...	REJ	RSTO	RSTOS0	RSTF
0	0.0	2.610418e-07	0.001057	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
1	0.0	3.446905e-07	0.000094	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
2	0.0	3.389216e-07	0.000259	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
3	0.0	3.158461e-07	0.000259	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
4	0.0	3.129617e-07	0.000394	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
...
494015	0.0	4.470881e-07	0.000365	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
494016	0.0	4.067060e-07	0.000443	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
494017	0.0	2.927706e-07	0.000233	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
494018	0.0	4.196859e-07	0.000233	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
494019	0.0	3.158461e-07	0.000239	0.0	0.0	0.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0

494020 rows x 119 columns

Splitting into training and testing data

```
In [19]: target=df1[numeric_cols].label
df2=df1[numeric_cols].drop('label',axis='columns')
```

```
In [20]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train, y_test=train_test_split(df2,target,test_size=.2)
```

Fitting data

```
In [21]: from sklearn.linear_model import LogisticRegression
```

```
In [22]: model = LogisticRegression()
```

```
In [23]: model.fit(X_train,y_train)
```

Out[23]: LogisticRegression()

```
In [24]: print(model.n_iter_[0])
```

94

```
In [25]: model.score(X_train,y_train)
```

Out[25]: 0.9997874579976519


```
In [26]: model.score(X_test,y_test)
```

Out[26]: 0.9997672159021902

```
In [27]: coefficients = model.coef_[0]
```

```
In [28]: scaled_weights=scaler.fit_transform(coefficients.reshape(118,1))
```

```
In [29]: fig = px.bar(
    x=coefficients,
    y=X_train.columns,
    orientation='h'
)
fig.show()
```



```
In [30]: model2_1 = LogisticRegression(solver='liblinear',penalty='l2',max_iter=12,C=0.01)
```

```
In [31]: model2_2 = LogisticRegression(solver='liblinear',penalty='l2',max_iter=20,C=10)
```

```
In [32]: model2_1.fit(X_train,y_train)
```

```
Out[32]: LogisticRegression(C=0.01, max_iter=12, solver='liblinear')
```

```
In [33]: model2_2.fit(X_train,y_train)
```

```
Out[33]: LogisticRegression(C=10, max_iter=20, solver='liblinear')
```

Training accuracy

```
In [34]: model2_1.score(X_train,y_train)
```

```
Out[34]: 0.998995486012712
```

```
In [35]: model2_2.score(X_train,y_train)
```

```
Out[35]: 0.9997823974737865
```

Testing accuracy

```
In [36]: model2_1.score(X_test,y_test)
```

```
Out[36]: 0.9987146269381807
```

```
In [37]: model2_2.score(X_test,y_test)
```

```
Out[37]: 0.9997570948544593
```

```
In [38]: coefficients = model2_1.coef_[0]
```

```
In [39]: scaled_weights=scaler.fit_transform(coefficients.reshape(118,1))
```

```
In [40]: fig = px.bar(
    x=coefficients,
    y=X_train.columns,
    orientation='h'
)
fig.show()
```

>

